


UNIT – 1

MACHINE LEARNING BASICS

Machine Learning Basics

Machine learning is a branch of artificial intelligence that enables algorithms to uncover hidden patterns within datasets. It allows them to predict new, similar data without explicit programming for each task. Machine learning finds applications in diverse fields such as image and speech recognition, natural language processing, recommendation systems, fraud detection, portfolio optimization, and automating tasks.


Machine learning's impact extends to autonomous vehicles, drones, and robots, enhancing their adaptability in dynamic environments. This approach marks a breakthrough where machines learn from data examples to generate accurate outcomes, closely intertwined with data mining and data science.



How does Machine Learning Work?

The process of machine learning consists of several basic steps that can be easily illustrated:

How machine learning work?




1. **Input Data:** The journey begins with data collection, encompassing various forms like photos, texts, sounds, and numerical figures.
2. **Data Preparation:** Once gathered, the data undergoes a cleaning process to ensure its accuracy. This phase involves rectifying inaccuracies, addressing errors, and filling in missing values. While this step is optional, it's highly recommended to achieve superior outcomes.
3. **Data Analysis:** At this juncture, the machine learning model delves into understanding the data's structure and discerning relationships between its various features.
4. **Pattern Search:** Leveraging insights from the analysis, the model unearths hidden patterns and connections within the data.
5. **Prediction:** Drawing on these identified patterns, the model is primed to forecast outcomes for new, unseen data.
6. **Decision-making:** Informed by the predictions, the system arrives at specific conclusions and takes corresponding actions. An example would be the system classifying an email as spam or legitimate; depending on this classification, it'll then allocate the email to the relevant folder.

Types of Machine Learning

There are several types of machine learning, each with special characteristics and applications. Some of the main types of machine learning algorithms are as follows:

1. Supervised Machine Learning
2. Unsupervised Machine Learning
3. Reinforcement Learning


Additionally, there is a more specific category called semi-supervised learning, which combines elements of both supervised and unsupervised learning.



1. Supervised Machine Learning

[Supervised learning](#) is defined as when a model gets trained on a "**Labelled Dataset**". Labelled datasets have both input and output parameters. In **Supervised Learning** algorithms learn to map points between inputs and correct outputs. It has both training and validation datasets labelled.

Supervised Learning




Example: Consider a scenario where you have to build an image classifier to differentiate between cats and dogs. If you feed the datasets of dogs and cats labelled images to the algorithm, the machine will learn to classify between a dog or a cat from these labeled images. When we input new dog or cat images that it has never seen before, it will use the learned algorithms and predict whether it is a dog or a cat. This is how **supervised learning** works, and this is particularly an image classification.

There are two main categories of supervised learning that are mentioned below:

- [Classification](#)
- [Regression](#)

Classification

Classification teaches a machine to sort things into categories. It learns by looking at examples with labels (like emails marked "spam" or "not spam"). After learning, it can decide which category new items belong to, like identifying if a new email is spam or not. For example a classification model might be trained on dataset of images labeled as either dogs or cats and it can be used to predict the class of new and unseen images as dogs or cats based on their features such as color, texture and shape.



Explaining classification in ml, **horizontal axis represents the combined values of color and texture features**. Vertical axis represents the combined values of shape and size features.

- Each **colored dot** in the plot represents an individual image, with the color indicating whether the model predicts the image to be a dog or a cat.
- The **shaded areas** in the plot show the **decision boundary**, which is the line or region that the model uses to decide which category (dog or cat) an image belongs to. The model classifies images on one side of the boundary as dogs and on the other side as cats, based on their features.

Types of Classification


When we talk about classification in machine learning, we're talking about the process of sorting data into categories based on specific features or characteristics. There are different types of classification problems depending on how many categories (or classes) we are working with and how they are organized. There are two main classification types in machine learning:

1. Binary Classification

- This is the simplest kind of classification. In binary classification, the goal is to sort the data into **two distinct categories**. Think of it like a simple choice between two options. Imagine a system that sorts emails into either **spam** or **not spam**. It works by looking at **different features of the email** like certain keywords or sender details, and decides whether it's spam or not. It only chooses between these two options.

2. Multiclass Classification

- Here, instead of just two categories, the data needs to be sorted into **more than two categories**. The model picks the one that best matches the input. Think of an image recognition system that sorts pictures of animals into categories like **cat**, **dog**, and **bird**.
- Basically, machine looks at the **features in the image (like shape, color, or texture)** and chooses which animal the picture is most likely to be based on the training it received.



3. Multi-Label Classification

In [multi-label classification](#) single piece of data can belong to **multiple categories** at once. Unlike multiclass classification where each data point belongs to only one class, multi-label classification allows **datapoints to belong to multiple classes**. A movie recommendation system could tag a movie as both **action** and **comedy**. The system checks various features (like movie plot, actors, or genre tags) and assigns multiple labels to a single piece of data, rather than just one.

Multilabel classification is relevant in specific use cases, but not as crucial for a starting overview of classification.


How does Classification in Machine Learning Work?

Classification involves training a model using a labeled dataset, where each input is paired with its correct output label. The model learns patterns and relationships in the data, so it can later predict labels for new, unseen inputs.

In machine learning, **classification** works by training a model to **learn patterns** from labeled data, so it can predict the category or class of new, unseen data. Here's how it works:

1. **Data Collection:** You start with a dataset where each item is labeled with the correct class (for example, "cat" or "dog").

2. **Feature Extraction:** The system identifies features (like color, shape, or texture) that help distinguish one class from another. These features are what the model uses to make predictions.
3. **Model Training:** Classification - machine learning algorithm uses the labeled data to learn how to map the features to the correct class. It looks for patterns and relationships in the data.
4. **Model Evaluation:** Once the model is trained, it's tested on new, unseen data to check how accurately it can classify the items.
5. **Prediction:** After being trained and evaluated, the model can be used to predict the class of new data based on the features it has learned.
6. **Model Evaluation:** Evaluating a classification model is a key step in machine learning. It helps us check how well the model performs and how good it is at handling new, unseen data. Depending on the problem and needs we can use different metrics to measure its performance.



If the quality metric is not satisfactory, the ML algorithm or hyperparameters can be adjusted, and the model is retrained. This iterative process continues until a satisfactory performance is achieved. In short, classification in machine learning is all about using existing labeled data to teach the model how to predict the class of new, unlabeled data based on the patterns it has learned.

Examples of Machine Learning Classification in Real Life

Classification algorithms are widely used in many real-world applications across various domains, including:

- **Email spam filtering**
- **Credit risk assessment:** Algorithms predict whether a loan applicant is likely to default by analyzing factors such as credit score, income, and loan history. This helps banks make informed lending decisions and minimize financial risk.
- **Medical diagnosis :** Machine learning models classify whether a patient has a certain condition (e.g., cancer or diabetes) based on medical data such as test results, symptoms, and patient history. This aids doctors in making quicker, more accurate diagnoses, improving patient care.
- **Image classification :** Applied in fields such as facial recognition, autonomous driving, and medical imaging.
- **Sentiment analysis:** Determining whether the sentiment of a piece of text is positive, negative, or neutral. Businesses use this to understand customer opinions, helping to improve products and services.
- **Fraud detection :** Algorithms detect fraudulent activities by analyzing transaction patterns and identifying anomalies crucial in protecting against credit card fraud and other financial crimes.
- **Recommendation systems :** Used to recommend products or content based on past user behavior, such as suggesting movies on Netflix or products on Amazon. This personalization boosts user satisfaction and sales for businesses.

Classification deals with predicting **categorical** target variables, which represent discrete classes or labels. For instance, classifying emails as spam or not spam, or predicting whether a patient has a high risk of heart disease. Classification algorithms learn to map the input features to one of the predefined classes.

Here are some classification algorithms:

- [Logistic Regression](#)
- [Support Vector Machine](#)
- [Random Forest](#)
- [Decision Tree](#)
- [K-Nearest Neighbors \(KNN\)](#)
- [Naive Bayes](#)

Regression

Regression, on the other hand, deals with predicting **continuous** target variables, which represent numerical values. For example, predicting the price of a house based on its size, location, and amenities, or forecasting the sales of a product. Regression algorithms learn to map the input features to a continuous numerical value.

Here are some regression algorithms:

- [Linear Regression](#)
- [Polynomial Regression](#)
- [Ridge Regression](#)
- [Lasso Regression](#)
- [Decision tree](#)
- [Random Forest](#)

Advantages of Supervised Machine Learning

- **Supervised Learning** models can have high accuracy as they are trained on **labelled** data.
- The process of decision-making in supervised learning models is often interpretable.
- It can often be used in pre-trained models which saves time and resources when developing new models from scratch.

Disadvantages of Supervised Machine Learning

- It has limitations in knowing patterns and may struggle with unseen or unexpected patterns that are not present in the training data.
- It can be time-consuming and costly as it relies on **labeled** data only.
- It may lead to poor generalizations based on new data.

Applications of Supervised Learning


Supervised learning is used in a wide variety of applications, including:

- **Image classification:** Identify objects, faces, and other features in images.
- **Natural language processing:** Extract information from text, such as sentiment, entities, and relationships.
- **Speech recognition:** Convert spoken language into text.
- **Recommendation systems:** Make personalized recommendations to users.
- **Predictive analytics:** Predict outcomes, such as sales, customer churn, and stock prices.

- **Medical diagnosis:** Detect diseases and other medical conditions.
- **Fraud detection:** Identify fraudulent transactions.
- **Autonomous vehicles:** Recognize and respond to objects in the environment.
- **Email spam detection:** Classify emails as spam or not spam.
- **Quality control in manufacturing:** Inspect products for defects.
- **Credit scoring:** Assess the risk of a borrower defaulting on a loan.
- **Gaming:** Recognize characters, analyze player behavior, and create NPCs.
- **Customer support:** Automate customer support tasks.
- **Weather forecasting:** Make predictions for temperature, precipitation, and other meteorological parameters.
- **Sports analytics:** Analyze player performance, make game predictions, and optimize strategies.

2. Unsupervised Machine Learning

[Unsupervised Learning](#) Unsupervised learning is a type of machine learning technique in which an algorithm discovers patterns and relationships using unlabeled data. Unlike supervised learning, unsupervised learning doesn't involve providing the algorithm with labeled target outputs. The primary goal of Unsupervised learning is often to discover hidden patterns, similarities, or clusters within the data, which can then be used for various purposes, such as data exploration, visualization, dimensionality reduction, and more.



Example: Consider that you have a dataset that contains information about the purchases you made from the shop. Through clustering, the algorithm can group the same purchasing behavior among you and other customers, which reveals potential customers without predefined labels. This type of information can help businesses get target customers as well as identify outliers.

There are two main categories of unsupervised learning that are mentioned below:

- [Clustering](#)
- [Association](#)

Clustering

Clustering is the process of grouping data points into clusters based on their similarity. This technique is useful for identifying patterns and relationships in data without the need for labeled examples.

Here are some clustering algorithms:


- [K-Means Clustering algorithm](#)
- [Mean-shift algorithm](#)
- [DBSCAN Algorithm](#)
- [Principal Component Analysis](#)
- [Independent Component Analysis](#)

What is Clustering?

The task of **grouping data points based on their similarity with each other** is called **Clustering or Cluster Analysis**. This method is defined under the branch of [unsupervised learning](#), which aims at gaining insights from unlabelled data points.


Think of it as you have a dataset of customers shopping habits. **Clustering can help you group customers with similar purchasing behaviors, which can then be used for targeted marketing, product recommendations, or customer segmentation**

For Example, In the graph given below, we can clearly see that there are 3 circular clusters forming on the basis of distance.



Now it is not necessary that the clusters formed **must be circular in shape**. The shape of clusters can be arbitrary. There are many algorithms that work well with detecting arbitrary shaped clusters.

For example, In the below given graph we can see that the clusters formed are not circular in shape.



Types of Clustering

Broadly speaking, there are 2 types of clustering that can be performed to group similar data points:

- **Hard Clustering:** In this type of clustering, each data point belongs to a cluster completely or not. For example, Let's say there are 4 data point and we have to cluster them into 2 clusters. So each data point will either belong to cluster 1 or cluster 2.

Data Points	Clusters
A	C1
B	C2
C	C2
D	C1

- **Soft Clustering:** In this type of clustering, instead of assigning each data point into a separate cluster, a probability or likelihood of that point being that cluster is evaluated. For example, Let's say there are 4 data point and we have to cluster them into 2 clusters. So we will be evaluating a probability of a data point belonging to both clusters. This probability is calculated for all data points.

Data Points	Probability of C1	Probability of C2
A	0.91	0.09
B	0.3	0.7
C	0.17	0.83
D	1	0

Uses of Clustering

Now before we begin with types of clustering algorithms, we will go through the use cases of Clustering algorithms. Clustering algorithms are majorly used for:

- **Market Segmentation:** Businesses use clustering to group their customers and use targeted advertisements to attract more audience.
- **Market Basket Analysis:** Shop owners analyze their sales and figure out which items are majorly bought together by the customers. For example, In USA, according to a study diapers and beers were usually bought together by fathers.
- **Social Network Analysis:** Social media sites use your data to understand your browsing behavior and provide you with targeted friend recommendations or content recommendations.
- **Medical Imaging:** Doctors use Clustering to find out diseased areas in diagnostic images like X-rays.
- **Anomaly Detection:** To find outliers in a stream of real-time dataset or forecasting fraudulent transactions we can use clustering to identify them.
- **Simplify working with large datasets:** Each cluster is given a cluster ID after clustering is complete. Now, you may reduce a feature set's whole feature set into its cluster ID. Clustering is effective when it can represent a complicated case with a straightforward cluster ID. Using the same principle, clustering data can make complex datasets simpler.

There are many more use cases for clustering but there are some of the major and common use cases of clustering. Moving forward we will be discussing Clustering Algorithms that will help you perform the above tasks.

Types of Clustering Methods

At the surface level, **clustering helps in the analysis of unstructured data. Graphing, the shortest distance, and the density of the data points are a few of the elements that influence cluster formation.** Clustering is the process of determining how related the objects are **based on a metric called the similarity measure.**

Similarity metrics **are easier to locate in smaller sets of features and harder as the number of features increases.** Depending on the type of clustering algorithm being utilized, several techniques are employed to group the data from the datasets. In this part, the clustering techniques are described. Various types of clustering algorithms are:

1. Centroid-based Clustering (Partitioning methods)
2. Density-based Clustering (Model-based methods)
3. Connectivity-based Clustering (Hierarchical clustering)

We will be going through each of these types in brief.

1. Centroid-based Clustering (Partitioning methods)

Centroid-based clustering organizes data points around central vectors (centroids) that represent clusters. Each data point belongs to the cluster with the nearest centroid. Generally, the similarity measure chosen for these algorithms are Euclidian distance, Manhattan Distance or Minkowski Distance.

The datasets are separated into a **predetermined number of clusters, and each cluster is referenced by a vector of values. When compared to the vector value, the input data variable shows no difference and joins the cluster.**

The major drawback for centroid-based algorithms is the requirement that we establish the number of clusters, "k," either intuitively or scientifically (using the Elbow Method) before any clustering machine learning system starts allocating the data points. Despite this limitation, it remains the most popular type of clustering due to its simplicity and efficiency. Popular algorithms of [Centroid-based clustering](#) are:

- [K-means](#) and
- [K-medoids](#) clustering

are some examples of this type clustering.

2. Density-based Clustering (Model-based methods)

Density-based clustering identifies clusters as areas of high density separated by regions of low density in the data space. Unlike centroid-based methods, density-based

clustering **automatically determines the number of clusters and is less susceptible to initialization positions.** **Key Characteristics:**

- Can find arbitrarily shaped clusters
- Handles noise and outliers well
- Excels with clusters of different sizes and shapes
- Ideal for datasets with irregularly shaped or overlapping clusters
- Effectively manages both dense and sparse data regions
- Focus on local density allows detection of various cluster morphologies

The most popular [density-based clustering](#) algorithm is [DBSCAN](#) and [OPTICS \(Ordering Points To Identify Clustering Structure\)](#).

3. Connectivity-based Clustering (Hierarchical clustering)

Connectivity-based clustering builds a **hierarchy of clusters using a measure of connectivity based on distance** when organizing a collection of items based on their similarities. This method builds a **dendrogram**, a tree-like structure that visually represents the relationships between objects.

At the base of the tree, each object starts as its own individual cluster. The algorithm then evaluates how similar the objects are to one another and begins merging the closest pairs of clusters into larger groups. This process continues iteratively, with clusters being combined step by step, until all objects are united into a single cluster at the top of the tree.

Association

[Association rule learning](#) is a technique for discovering relationships between items in a dataset. It identifies rules that indicate the presence of one item implies the presence of another item with a specific probability.

Here are some association rule learning algorithms:

- [Apriori Algorithm](#)
- [Eclat](#)
- [FP-growth Algorithm](#)

Association rule mining finds interesting associations and relationships among large sets of data items. This rule shows how frequently a itemset occurs in a transaction. A typical example is a Market Based Analysis. Market Based Analysis is one of the key techniques used by large relations to show associations between items. It allows retailers to identify relationships between the items that people buy together frequently. Given a set of transactions, we can find rules that will predict the occurrence of an item based on the occurrences of other items in the transaction.

TID	Items
1	Bread, Milk
2	Bread, Diaper, Beer, Eggs
3	Milk, Diaper, Beer, Coke
4	Bread, Milk, Diaper, Beer
5	Bread, Milk, Diaper, Coke

Before we start defining the rule, let us first see the basic definitions. **Support Count($\sigma\sigma$)** - Frequency of occurrence of a itemset.

Here $\sigma\sigma(\{\text{Milk, Bread, Diaper}\})=2$

Frequent Itemset - An itemset whose support is greater than or equal to minsup threshold. **Association Rule** - An implication expression of the form $X \rightarrow Y$, where X and Y are any 2 itemsets.

Example: $\{\text{Milk, Diaper}\} \rightarrow \{\text{Beer}\}$

Rule Evaluation Metrics -

- **Support(s)** - The number of transactions that include items in the $\{X\}$ and $\{Y\}$ parts of the rule as a percentage of the total number of transaction. It is a measure of how frequently the collection of items occur together as a percentage of all transactions.
- **Support = $\sigma\sigma(X+Y) \div \text{total}$** - It is interpreted as fraction of transactions that contain both X and Y.
- **Confidence(c)** - It is the ratio of the no of transactions that includes all items in $\{B\}$ as well as the no of transactions that includes all items in $\{A\}$ to the no of transactions that includes all items in $\{A\}$.
- **Conf($X \Rightarrow Y$) = $\text{Supp}(X \cup Y) \div \text{Supp}(X)$** - It measures how often each item in Y appears in transactions that contains items in X also.
- **Lift(l)** - The lift of the rule $X \Rightarrow Y$ is the confidence of the rule divided by the expected confidence, assuming that the itemsets X and Y are independent of each other. The expected confidence is the confidence divided by the frequency of $\{Y\}$.

- **Lift(X=>Y) = Conf(X=>Y) ÷÷ Supp(Y)** - Lift value near 1 indicates X and Y almost often appear together as expected, greater than 1 means they appear together more than expected and less than 1 means they appear less than expected. Greater lift values indicate stronger association.

Example - From the above table, {Milk, Diaper}=>{Beer}

$$s = \sigma(\{\text{Milk, Diaper, Beer}\}) \div |T|$$

$$= 2/5$$

$$= 0.4$$

$$c = \sigma(\text{Milk, Diaper, Beer}) \div \sigma(\text{Milk, Diaper})$$

$$= 2/3$$

$$= 0.67$$

$$l = \text{Supp}(\{\text{Milk, Diaper, Beer}\}) \div \text{Supp}(\{\text{Milk, Diaper}\}) * \text{Supp}(\{\text{Beer}\})$$

$$= 0.4 / (0.6 * 0.6)$$

$$= 1.11$$

The Association rule is very useful in analyzing datasets. The data is collected using bar-code scanners in supermarkets. Such databases consists of a large number of transaction records which list all items bought by a customer on a single purchase. So the manager could know if certain groups of items are consistently purchased together and use this data for adjusting store layouts, cross-selling, promotions based on statistics.

Advantages of Unsupervised Machine Learning

- It helps to discover hidden patterns and various relationships between the data.
- Used for tasks such as **customer segmentation, anomaly detection, and data exploration.**
- It does not require labeled data and reduces the effort of data labeling.

Disadvantages of Unsupervised Machine Learning

- Without using labels, it may be difficult to predict the quality of the model's output.
- Cluster Interpretability may not be clear and may not have meaningful interpretations.
- It has techniques such as [autoencoders](#) and [dimensionality reduction](#) that can be used to extract meaningful features from raw data.

Applications of Unsupervised Learning

Here are some common applications of unsupervised learning:


- **Clustering:** Group similar data points into clusters.
- **Anomaly detection:** Identify outliers or anomalies in data.
- **Dimensionality reduction:** Reduce the dimensionality of data while preserving its essential information.
- **Recommendation systems:** Suggest products, movies, or content to users based on their historical behavior or preferences.
- **Topic modeling:** Discover latent topics within a collection of documents.
- **Density estimation:** Estimate the probability density function of data.
- **Image and video compression:** Reduce the amount of storage required for multimedia content.
- **Data preprocessing:** Help with data preprocessing tasks such as data cleaning, imputation of missing values, and data scaling.
- **Market basket analysis:** Discover associations between products.
- **Genomic data analysis:** Identify patterns or group genes with similar expression profiles.
- **Image segmentation:** Segment images into meaningful regions.
- **Community detection in social networks:** Identify communities or groups of individuals with similar interests or connections.
- **Customer behavior analysis:** Uncover patterns and insights for better marketing and product recommendations.
- **Content recommendation:** Classify and tag content to make it easier to recommend similar items to users.
- **Exploratory data analysis (EDA):** Explore data and gain insights before defining specific tasks.

3. Reinforcement Machine Learning

[Reinforcement machine learning](#) algorithm is a learning method that interacts with the environment by producing actions and discovering errors. **Trial, error, and delay** are the most relevant characteristics of reinforcement learning. In this technique, the model keeps on increasing its performance using Reward Feedback to learn the behavior or pattern. These algorithms are specific to a particular problem e.g. Google Self Driving car, AlphaGo where a bot competes with humans and even itself to get better and better performers in Go Game. Each time we feed in data, they learn and add the data to their knowledge which is training data. So, the more it learns the better it gets trained and hence experienced.

Here are some of most common reinforcement learning algorithms:

- **Q-learning**: Q-learning is a model-free RL algorithm that learns a Q-function, which maps states to actions. The Q-function estimates the expected reward of taking a particular action in a given state.
- **SARSA (State-Action-Reward-State-Action)**: SARSA is another model-free RL algorithm that learns a Q-function. However, unlike Q-learning, SARSA updates the Q-function for the action that was actually taken, rather than the optimal action.
- **Deep Q-learning**: Deep Q-learning is a combination of Q-learning and deep learning. Deep Q-learning uses a neural network to represent the Q-function, which allows it to learn complex relationships between states and actions.



Example: Consider that you are training an [AI](#) agent to play a game like chess. The agent explores different moves and receives positive or negative feedback based on the outcome. Reinforcement Learning also finds applications in which they learn to perform tasks by interacting with their surroundings.

Types of Reinforcement Machine Learning

There are two main types of reinforcement learning:

Positive reinforcement

- Rewards the agent for taking a desired action.
- Encourages the agent to repeat the behavior.
- Examples: Giving a treat to a dog for sitting, providing a point in a game for a correct answer.

Negative reinforcement

- Removes an undesirable stimulus to encourage a desired behavior.
- Discourages the agent from repeating the behavior.

- Examples: Turning off a loud buzzer when a lever is pressed, avoiding a penalty by completing a task.

Advantages of Reinforcement Machine Learning

- It has autonomous decision-making that is well-suited for tasks and that can learn to make a sequence of decisions, like robotics and game-playing.
- This technique is preferred to achieve long-term results that are very difficult to achieve.
- It is used to solve a complex problems that cannot be solved by conventional techniques.

Disadvantages of Reinforcement Machine Learning

- Training Reinforcement Learning agents can be computationally expensive and time-consuming.
- Reinforcement learning is not preferable to solving simple problems.
- It needs a lot of data and a lot of computation, which makes it impractical and costly.

Applications of Reinforcement Machine Learning

Here are some applications of reinforcement learning:


- **Game Playing:** RL can teach agents to play games, even complex ones.
- **Robotics:** RL can teach robots to perform tasks autonomously.
- **Autonomous Vehicles:** RL can help self-driving cars navigate and make decisions.
- **Recommendation Systems:** RL can enhance recommendation algorithms by learning user preferences.
- **Healthcare:** RL can be used to optimize treatment plans and drug discovery.
- **Natural Language Processing (NLP):** RL can be used in dialogue systems and chatbots.
- **Finance and Trading:** RL can be used for algorithmic trading.
- **Supply Chain and Inventory Management:** RL can be used to optimize supply chain operations.
- **Energy Management:** RL can be used to optimize energy consumption.
- **Game AI:** RL can be used to create more intelligent and adaptive NPCs in video games.
- **Adaptive Personal Assistants:** RL can be used to improve personal assistants.
- **Virtual Reality (VR) and Augmented Reality (AR):** RL can be used to create immersive and interactive experiences.
- **Industrial Control:** RL can be used to optimize industrial processes.

- **Education:** RL can be used to create adaptive learning systems.
- **Agriculture:** RL can be used to optimize agricultural operations.

Semi-Supervised Learning: Supervised + Unsupervised Learning

Semi-Supervised learning is a machine learning algorithm that works between the supervised and unsupervised learning so it uses both **labelled and unlabelled** data. It's particularly useful when obtaining labeled data is costly, time-consuming, or resource-intensive. This approach is useful when the dataset is expensive and time-consuming. Semi-supervised learning is chosen when labeled data requires skills and relevant resources in order to train or learn from it.

We use these techniques when we are dealing with data that is a little bit labeled and the rest large portion of it is unlabeled. We can use the unsupervised techniques to predict labels and then feed these labels to supervised techniques. This technique is mostly applicable in the case of image data sets where usually all images are not labeled.



Example: Consider that we are building a language translation model, having labeled translations for every sentence pair can be resources intensive. It allows the models to learn from labeled and unlabeled sentence pairs, making them more accurate. This technique has led to significant improvements in the quality of machine translation services.

Types of Semi-Supervised Learning Methods

There are a number of different semi-supervised learning methods each with its own characteristics. Some of the most common ones include:

- **Graph-based semi-supervised learning:** This approach uses a graph to represent the relationships between the data points. The graph is then used to propagate labels from the labeled data points to the unlabeled data points.
- **Label propagation:** This approach iteratively propagates labels from the labeled data points to the unlabeled data points, based on the similarities between the data points.

- **Co-training:** This approach trains two different machine learning models on different subsets of the unlabeled data. The two models are then used to label each other's predictions.
- **Self-training:** This approach trains a machine learning model on the labeled data and then uses the model to predict labels for the unlabeled data. The model is then retrained on the labeled data and the predicted labels for the unlabeled data.
- **Generative adversarial networks (GANs):** GANs are a type of deep learning algorithm that can be used to generate synthetic data. GANs can be used to generate unlabeled data for semi-supervised learning by training two neural networks, a generator and a discriminator.

Advantages of Semi- Supervised Machine Learning

- It leads to better generalization as compared to **supervised learning**, as it takes both labeled and unlabeled data.
- Can be applied to a wide range of data.

Disadvantages of Semi- Supervised Machine Learning

- **Semi-supervised** methods can be more complex to implement compared to other approaches.
- It still requires some **labeled data** that might not always be available or easy to obtain.
- The unlabeled data can impact the model performance accordingly.

Applications of Semi-Supervised Learning

Here are some common applications of semi-supervised learning:


- **Image Classification and Object Recognition:** Improve the accuracy of models by combining a small set of labeled images with a larger set of unlabeled images.
- **Natural Language Processing (NLP):** Enhance the performance of language models and classifiers by combining a small set of labeled text data with a vast amount of unlabeled text.
- **Speech Recognition:** Improve the accuracy of speech recognition by leveraging a limited amount of transcribed speech data and a more extensive set of unlabeled audio.
- **Recommendation Systems:** Improve the accuracy of personalized recommendations by supplementing a sparse set of user-item interactions (labeled data) with a wealth of unlabeled user behavior data.
- **Healthcare and Medical Imaging:** Enhance medical image analysis by utilizing a small set of labeled medical images alongside a larger set of unlabeled images.

Supervised vs Unsupervised vs Reinforcement Learning

Machine learning (ML) is a subset of artificial intelligence (AI). It enables systems to learn from data, identify patterns and make decisions with minimal human intervention. The three primary types of ML are:

1. **Supervised Learning:** Learning from labelled data.
2. **Unsupervised Learning:** Discovering patterns in unlabeled data.
3. **Reinforcement Learning:** Learning through interactions with an environment.

Each approach has unique characteristics, advantages and real-world applications.



Supervised Learning

Supervised learning is like learning with a teacher. The model is trained on a labeled dataset, meaning each input has a corresponding output. The key characteristics of supervised learning are:

- **Labeled Data:** Training data has predefined labels.
- **Types of Problems:** Used for **classification task** like spam detection and **regression task** like predicting house prices.
- **Algorithms:** Linear Regression, Logistic Regression, SVM, Decision Trees, Neural Networks.

Unsupervised Learning

Unsupervised learning works with data that has no predefined labels. The model identifies patterns, clusters or associations independently. The key characteristics of unsupervised learning are:

- **Unlabeled Data:** No predefined outputs.
- **Types of Problems:** Used for **Clustering task** like customer segmentation and **association task** like market basket analysis.
- **Algorithms:** K-Means, Hierarchical Clustering, PCA, Autoencoders.

Reinforcement Learning (RL)

Reinforcement learning involves an agent that interacts with an environment, learning through rewards and penalties to maximize long-term success. The key characteristics of reinforcement learning are:

- **Interaction-Based Learning:** The agent learns by taking actions and receiving feedback.
- **No Labeled Data:** Learns from trial and error.
- **Algorithms:** Q-learning, SARSA, Deep Q-Networks (DQN).

Comparison Table: Supervised vs Unsupervised vs Reinforcement Learning

Criteria	Supervised Learning	Unsupervised Learning	Reinforcement Learning
Definition	Learns from labeled data	Identifies patterns in unlabeled data	Learns through interaction with environment
Type of Data	Labeled data	Unlabeled data	No predefined data learn from environment
Type of Problems	Classification, Regression	Clustering, Association	Sequential decision-making
Supervision	Requires external supervision	No supervision	No supervision, learns from feedback

Criteria	Supervised Learning		Unsupervised Learning		Reinforcement Learning	
Algorithms	SVM, Trees, Networks	Decision Neural	K-Means, Autoencoders	PCA,	Q-learning, SARSA	DQN,
Goal	Predict accurately	outcomes	Discover patterns	hidden	Optimize actions for maximum rewards	
Applications	Medical diagnosis, fraud detection		Customer segmentation, anomaly detection		Self-driving cars, robotics, gaming	


Real-World Applications

Machine Learning Type	Domain	Examples
Supervised Learning	Healthcare	Disease diagnosis like cancer detection
	Finance	Loan approval, credit risk assessment
	NLP	Sentiment analysis, text classification
Unsupervised Learning	E-commerce	Product recommendation, customer segmentation
	Cybersecurity	Fraud detection, intrusion detection
	Biology	Gene classification, dimensionality reduction

Machine Learning Type	Domain	Examples
Reinforcement Learning	Autonomous Driving	Self-driving cars learning optimal driving behavior
	Robotics	Training robots for automated assembly tasks
	Gaming	AI-driven strategy games like AlphaGo

Machine Learning Models

Machine Learning models are very powerful resources that automate multiple tasks and make them more accurate and efficient. ML handles new data and scales the growing demand for technology with valuable insight. It improves the performance over time. This cutting-edge technology has various benefits such as faster processing or response, enhancement of decision-making, and specialized services. In this article, we will discuss ***Machine Learning Models, their types, How Machine Learning works, Real-world examples of ML Models, and the Future of Machine Learning Models.***



A model of machine learning is a set of programs that can be used to find the pattern and make a decision from an unseen dataset. These days [NLP](#) (Natural language Processing) uses the machine learning model to recognize the unstructured text into usable data and insights. You may have heard about image recognition which is used to identify objects such as boy, girl, mirror, car, dog, etc. A model always requires a dataset to perform various tasks during training.

In training duration, we use a [machine learning algorithm](#) for the optimization process to find certain patterns or outputs from the [dataset](#) based upon tasks.

Types of Machine Learning Models

Machine learning models can be broadly categorized into four main paradigms based on the type of data and learning goals:

1. Supervised Models

Supervised learning is the study of algorithms that use labeled data in which each data instance has a known category or value to which it belongs. This results in the model to discover the relationship between the input features and the target outcome.

1.1 Classification

The classifier algorithms are designed to indicate whether a new data point belongs to one or another among several predefined classes. Imagine when you are organising emails into spam or inbox, categorising images as cat or dog, or predicting whether a loan applicant is a credible borrower. In the classification models, there is a learning process by the use of labeled examples from each category. In this process, they discover the correlations and relations within the data that help to distinguish class one from the other classes. After learning these patterns, the model is then capable of assigning these class labels to unseen data points.

Common Classification Algorithms:

- **Logistic Regression:** A very efficient technique for the classification problems of binary nature (two types, for example, spam/not spam).
- **[Support Vector Machine \(SVM\)](#):** Good for tasks like classification, especially when the data has a large number of features.
- **Decision Tree:** Constructs a decision tree having branches and proceeds to the class predictions through features.
- **[Random Forest](#):** The model generates an "ensemble" of decision trees that ultimately raise the accuracy and avoid overfitting (meaning that the model performs great on the training data but lousily on unseen data).
- **[K-Nearest Neighbors \(KNN\)](#):** Assigns a label of the nearest neighbors for a given data point.

1.2 Regression

Regression algorithms are about forecasting of a continuous output variable using the input features as their basis. This value could be anything such as predicting real estate prices or stock market trends to anticipating customer churn (how likely customers stay) and sales forecasting. Regression models make the use of features to understand the relationship among the continuous features and the output variable. That is, they use the pattern that is learned to determine the value of the new data points.

Common Regression Algorithms

- **Linear Regression:** Fits depth of a line to the data to model for the relationship between features and the continuous output.
- **Polynomial Regression:** Similiar to linear regression but uses more complex polynomial functions such as quadratic, cubic, etc, for accommodating non-linear relationships of the data.
- **Decision Tree Regression:** Implements a decision tree-based algorithm that predicts a continuous output variable from a number of branching decisions.
- **Random Forest Regression:** Creates one from several decision trees to guarantee error-free and robust regression prediction results.
- **Support Vector Regression (SVR):** Adjusts the Support Vector Machine ideas for regression tasks, where we are trying to find one hyperplane that most closely reflects continuous output data.

2. Unsupervised Models

Unsupervised learning involves a difficult task of working with data which is not provided with pre-defined categories or label.

2.1 Clustering

Visualize being given a basket of fruits with no labels on them. The fruits clustering algorithms are to group them according to the inbuilt similarities. Techniques like K-means clustering are defined by exact number of clusters ("red fruits" and "green fruits") and then each data point (fruit) is assigned to the cluster with the highest similarity within based on features (color, size, texture). Contrary to this, hierarchical clustering features construction of hierarchy of clusters which makes it more easy to study the system of groups. Spatial clustering algorithm Density-Based Spatial Clustering of Applications with Noise (DBSCAN) detects groups of high-density data points, even in those areas where there is a lack of data or outliers.

2.2 Dimensionality Reduction

Sometimes it is difficult to both visualize and analyze the data when you have a large feature space (dimensions). The purpose of dimensionality reduction methods is to decrease the dimensions needed to maintain the key features. Dimensions of greatest importance are identified by principal component analysis (PCA), which is the reason why data is concentrated in fewer dimensions with the highest variations. This speeds up model training as well as offers a chance for more efficient visualization. LDA (Linear Discriminant Analysis) also resembles PCA but it is made for classification tasks where it concentrates on dimensions that can differentiate the present classes in the dataset.

2.3 Anomaly Detection

Unsupervised learning can also be applied to find those data points which greatly differ than the majorities. The statistics model may identify these outliers, or anomalies as signaling of

errors, fraud or even something unusual. Local Outlier Factor (LOF) makes a comparison of a given data point's local density with those surrounding it. It then flags out the data points with significantly lower densities as outliers or potential anomalies. Isolation Forest is the one which uses different approach, which is to recursively isolate data points according to their features. Anomalies usually are simple to contemplate as they often necessitate fewer steps than an average normal point.

3. Semi-Supervised Model

Besides, supervised learning is such a kind of learning with labeled data that unsupervised learning, on the other hand, solves the task where there is no labeled data. Lastly, semi-supervised learning fills the gap between the two. It reveals the strengths of both approaches by training using data sets labeled along with unlabeled one. This is especially the case when labeled data might be sparse or prohibitively expensive to acquire, while unlabeled data is undoubtedly available in abundance.

3.1 Generative Semi-Supervised Learning

Envision having a few pictures of cats with labels and a universe of unlabeled photos. The big advantage of generative semi-supervised learning is its utilization of such a scenario. It exploits a generative model to investigate the unlabeled pictures and discover the orchestrating factors that characterize the data. This technique can then be used to generate the new synthetic data points that have the same features with the unlabeled data. The synthetic data is then labeled with the pseudo-labels that the generative model has interpreted from the data. This approach combines the existing labeled data with the newly generated labeled data to train the final model which is likely to perform better than the previous model that was trained with only the limited amount of the original labeled data.

3.2 Graph-based Semi-Supervised Learning

This process makes use of the relationships between data points and propagates labels to unmarked ones via labeled ones. Picture a social network platform where some of the users have been marked as fans of sports (labeled data). Cluster-based methods can analyze the links between users (friendships) and even apply this information to infer that if a user is connected to someone with a "sports" label then this user might also be interested in sports (unbiased labels with propagated label). While links and the entire structure of the network are also important for the distribution of labels. This method is beneficial when the data points are themselves connected to each other and this connection can be exploiting during labelling of new data.

4. Reinforcement learning Models

Reinforcement learning takes a dissimilar approach from [supervised learning](#) and unsupervised learning. Different from supervised learning or just plain discovery of hidden patterns, reinforcement learning adopt an agent as it interacts with the surrounding and learns. This agent is a learning one which develops via experiment and error, getting rewarded for the desired

actions and punished for the undesired ones. The main purpose is to help players play the game that can result in the highest rewards.

4.1 Value-based learning:

Visualize a robot trying to find its way through a maze. It has neither a map nor instructions, but it gets points for consuming the cheese at the end and fails with deduction of time when it runs into a wall. Value learning is an offshoot of predicting the anticipated future reward of taking a step in a particular state. For example, the algorithm Q-learning will learn a Q-value for each state-action combination. This Q-value is the expected reward for that action at that specific state. Through a repetitive process of assessing the state, gaining rewards, and updating the Q-values the agent manages to determine that which actions are most valuable in each state and eventually guides it to the most rewarding path. In contrast, [SARSA \(State-Action-Reward-State-Action\)](#) looks at the value of the succeeding state-action pair that influences the exploration strategy.

4.2 Policy-based learning:

In contrast to the value-based learning, where we are learning a specific value for each state-action pair, in policy-based learning we are trying to directly learn a policy which maps states to actions. This policy in essence commands the agent to act in different situations as specified by the way it is written. Actor-Critic is a common approach that combines two models: an actor that retrains the policy and a critic that retrains the value function (just like value-based methods). The actor witnesses the critic's feedback which updates the policy that the actor uses for better decision making. Proximal Policy Optimization (PPO) is a specific policy-based method which focuses on high variance issues that complicate early policy-based learning methods.

Applications of Machine Learning

Machine Learning (ML) is one of the most significant advancements in the field of technology. It gives machines the ability to learn from data and improve over time without being explicitly programmed. ML models identify patterns from data and use them to make predictions or decisions.

Organizations use machine learning to automate tasks, make smarter decisions and gain valuable insights. ML is shaping the world around us. Here are few real-world applications of Machine Learning:

1. Healthcare and Medical Diagnosis

ML algorithms can analyze large volumes of patient data, medical scans and genetic information to aid in diagnosis and treatment.

Applications:

- **Disease Detection:** ML models are used to identify diseases like cancer, pneumonia and Parkinson's from medical images. They often achieve accuracy comparable to or better than human doctors.

- **Predictive Analytics:** By analyzing patient history and symptoms, models can predict the risk of certain diseases or potential complications.
- **Drug Discovery:** ML accelerates the drug development process by predicting how different compounds will interact, reducing the time and cost of research.

2. Smart Assistants and Human-Machine Interaction

Virtual assistants systems rely on [natural language processing \(NLP\)](#) and [speech recognition](#) to understand commands and respond intelligently.

Applications:

- **Voice Assistants:** Tools like Siri, Alexa and Google Assistant convert spoken input into actionable commands.
- **Voice Search & Transcription:** ML enables users to perform hands-free web searches and get transcription during meetings or phone calls.
- **Chatbots:** Businesses use AI-powered chatbots for 24/7 customer support, helping resolve queries faster and more efficiently.

3. Personalized Recommendations and User Experience

Modern digital platforms uses personalization which is done by using [recommender systems](#). Machine learning models analyze user behavior to deliver relevant content, improving engagement and satisfaction.

Applications:

- **Streaming Platforms:** Netflix and Spotify suggest shows and songs based on your watching or listening history.
- **E-commerce:** Sites like Amazon recommend products tailored to your preferences, browsing patterns and past purchases.
- **Social Media:** Algorithms curate content feeds, prioritize posts and suggest friends or pages.

These systems use techniques like [collaborative filtering](#) and [content-based filtering](#) to create personalized digital experiences.

4. Fraud Detection and Financial Forecasting

In finance, vast sums of money move digitally and machine learning plays a important role in fraud detection and market analysis.

Applications:

- **Transaction Monitoring:** Banks use ML models to detect unusual spending behavior and flag suspicious transactions.

- **Loan Risk Assessment:** Credit scoring models analyze customer profiles and predict the likelihood of default.
- **Stock Market Prediction:** ML is used to analyze historical stock data and forecast price movements. Stock markets are complex, algorithmic trading uses these predictions for better decision-making.

5. Autonomous Vehicles and Smart Mobility

[Self-driving vehicles](#) use ML to understand their environment, navigate safely and make immediate decisions.

Key Components:

- **Computer Vision:** Recognizing lanes, pedestrians, traffic signals and obstacles.
- **Sensor Fusion:** Combining data from cameras, LiDAR and radar for a 360-degree view.
- **Behavior Prediction:** Anticipating how other drivers or pedestrians may act.

Autonomous vehicles are capable of operating with minimal human input. Beyond cars, ML is also being used in traffic optimization, smart navigation systems and predictive maintenance in transportation.